Sumedh Pendurkar

sumedhpendurkar@tamu.edu | sumedhpendurkar.github.io | sumedhpendurkar | Scholar

Research Interests / Research Experience

Reinforcement Learning, Large Language Models (LLMs), Foundation Models, LLM Alignment, LLM Agents, Combinatorial Optimization

Education

Texas A&M University — *TX, USA*

Doctor of Philosophy in Computer Science, 4/4 GPA — *August 2019 - August 2025*

- *Key courses*: Reinforcement Learning, Applied Bayes Methods, Optimization for Machine Learning, Machine Learning, AI, Analysis of Algorithms, Algorithms for Graph Mining
- *Thesis Title:* Learnable Guiding Functions for State-Space Search Algorithms
- *Advisor*: Dr. Guni Sharon

College of Engineering, Pune — *Pune, India*

Bachelor of Technology in Computer Engineering, 9.12/10 CGPA — *July 2015 - May 2019*

- *Key courses*: Data Science, Design and Analysis of Algorithms, AI, Theory of Computation, Introduction to Graph Theory

Work Experience

Meta Platforms — *Menlo Park, CA*

Research Scientist — *June 2025 - Present*

- Advanced core modeling techniques for foundation models for recommendation systems.
- Designed model architecture for foundation model for Facebook feed.
- *Technologies:* PyTorch, Triton, Recommendation Systems, Machine Learning, Foundation Models.

Decompute Inc. — *Remote*

AI Researcher Intern — *Jan 2025 - March 2025*

- Improving LLM training and inference computational performance across distributed systems.
- Evaluation of LoRA based methods for fine tuning LLMs on several datasets.
- *Technologies:* PyTorch, Apple MLX, Hugging Face, openMPI

University of Alberta — *Edmonton, AB, Canada*

Visiting Student with Dr. Nathan Sturtevant and Dr. Levi Lelis — *May 2023 - July 2023*

- Worked on developing curriculum generation method (TSC) for various guided state-space search algorithms.
- TSC achieved 5-36 times faster performance as compared to the baseline algorithms (Publication @ SoCS)

Niantic Inc. — *Sunnyvale, CA, USA*

Machine Learning Scientist Intern — *May 2022 - Aug 2022*

- Developed a novel formulation for game meta balance problem and a novel method (BiGMB) to optimize it.
- BiGMB scaled orders-of-magnitude better than previous methods for harder problems (Publication @ AAMAS)
- *Technologies:* PyTorch, OpenAI gym

Goldman Sachs — *Bangalore, India*

Summer Technology Analyst (Intern) — *May 2018 - July 2018*

- Worked on UI part of a change management tool & developed RESTful web services for change management tool.
- *Technologies:* Java, TypeScript, JavaScript, REST API

Indian Institute of Technology (IIT), Roorkee — *Roorkee, India*

Visiting Student with Dr. Biplab Banerjee — *May 2017 - July 2017*

- Designed deconv-net based model for single image super-resolution to enhance optical satellite images.
- Developed model achieved 0.55 dB PSNR over previous state-of-the-art. (Publication @ICIAP)
- Investigated zero-shot techniques for super-resolution of optical satellite images.

Selected Publications

Policy-Guided Search on Tree-of-Thoughts for Efficient Problem Solving with Bounded Language Model Queries
S. Pendurkar, G. Sharon — *2025*
Transactions of Machine Learning Research (TMLR)

Goal Distribution in Conflict-Based Search for Multi-Agent Pathfinding and its Implications to Monte-Carlo Sampling
S. Pendurkar, C. Simpson, S. Fayaz, G. Sharon — *2025*
AAAI workshop on Multi-Agent Path Finding (MAPF)

Computer Vision Approach for Analysis of Numerical and Experimental Detonation Cellular Structure Images
D. Jalontzki, A. Zussman, S. Pendurkar, G. Sharon, Y. Kozak — *2025*
Israel Annual Conference on Aerospace Sciences, IACAS

Curriculum Generation for Learning Guiding Functions in State-Space Search Algorithms — *AB, Canada*
S. Pendurkar, L. Lelis, N. Sturtevant, G. Sharon — *2024*
Symposium on Combinatorial Search (SoCS)

The (Un)Scalability of Informed Heuristic Function Estimation in NP-Hard Search Problems
S. Pendurkar, T. Huang, B. Juba, J.Zhang, S. Koenig, G. Sharon — *2023*
Transactions of Machine Learning Research (TMLR)

Bilevel Entropy based Mechanism Design for Balancing Meta in Video Games — *London, UK*
S. Pendurkar, C. Chow, J. Luo, G. Sharon — *2023*
International Conference on Autonomous Agents and Multiagent Systems (AAMAS)

Comparison between popular Genetic Algorithm (GA)-based tool and Covariance Matrix Adaptation - Evolutionary Strategy (CMA-ES) for optimizing indoor daylight — *China*
M. Anis, S. Pendurkar, Y. Yi, G. Sharon — *2023*
IBPSA International Conference and Exhibition on Building Simulation

The (Un)Scalability of Heuristic Approximators for NP-Hard Search Problems — *New Orleans, USA*
S. Pendurkar, T. Huang, S. Koenig, G. Sharon — *2022*
Proceedings of NeurIPS workshop. ICBINB.

A Discussion on the Scalability of Heuristic Approximators — *Vienna, Austria*
S. Pendurkar, T. Huang, S. Koenig, G. Sharon — *2022*
Symposium on Combinatorial Search (SoCS) (Extended Abstract)

A Joint Imitation-Reinforcement Learning Framework for Reduced Baseline Regret — *Prague, Czech Republic*
S. Dey, S. Pendurkar, G. Sharon, JP. Hanna — *2021*
International Conference on Intelligent Robots and Systems (IROS)

Single Image Super-Resolution for Optical Satellite Scenes Using Deep Deconvolutional Network — *Trento, Italy*
S. Pendurkar, B. Banerjee, S. Saha, F. Bovolo — *2019*
International Conference on Image Analysis and Processing (ICIAP)

Semantic Guided Deep Unsupervised Image Segmentation — *Trento, Italy*
S. Saha, B. Banerjee, S. Sudhakaran, S. Pendurkar — *2019*
International Conference on Image Analysis and Processing (ICIAP)

Technical Skills

Proficient: Programming	Python, C, JavaScript
Intermediate: Programming	Java, C++, SQL
Tools and Libraries	PyTorch, Keras, Apple MLX, Git, Angular, openMPI, GTK, Latex

Other Projects

Autograder for the Deep Reinforcement Learning Course *August 2022 - December 2023*

- Designed test cases and developed autograder for CSCE 642 course at Texas A&M University.
- The autograder is currently used by other universities.

Sampling an action from a Q function in continuous action spaces *August 2021 - May 2022*

- Investigated various sampling techniques, to efficiently sample actions from a neural network parameterized Q function.

Light-Regularized-GANs for low light images *September 2019 - Jan 2021*

- Added an intensity based regularisation to LightEnhancementGAN, to control the intensity of light added to the image without any external supervision.

Open-Ended Visual Question Answering System *April 2018 - May 2019*

- Designed an attention based multi-modal fusion model which gives a free flowing answer to a question based on video.

Word completion feature for GNU-Nano text editor *July 2016 - December 2016*

- Authored a word-completion feature which completes the current word based on the text present in the open file.
- This feature was incorporated in GNU-Nano, an open source project.

Communication/on-board controller system for pico satellite *April 2016 - July 2018*

- Developed shared memory protocols for two asynchronous controllers for on-board data sharing on a pico-satellite.
- Worked on interfacing various peripherals with on-board controllers for data collection.

Honors & Awards

2020	**First Place**, 2020 TAMIDS Data Science Competition	*TX, USA*
2018	**Deloitte Innovation Award**, Ministry of Road and Railways, Smart India Hackathon	*Nagpur, India*
2018	**Finalist (40/1980)**, Philips Hackathon on Data Science	*Bangalore, India*
2013	**Scholarship Holder**, National Talent Search Exam (NTSE), awarded to top 1000 students in India	*India*

Professional Activities

2020	**Reviewer**, ICRA 2021
2021	**Reviewer**, IROS 2021
2022	**Program Committee**, AAAI 2023, AAAI workshop on multi-agent path finding
2023	**Program Committee**, NeurIPS 2023, AAAI 2023, NeurIPS workshop 2023
2023	**Student Volunteer**, AAMAS 2023
2024	**Program Committee**, ICML 2024, TMLR 2024, AAAI 2024
2025	**Program Committee**, ICML 2025, AAAI 2026